

Automated my own day-to-day work as QA Lead on an \$11B AI legal platform: an AI agent creates, investigates and repairs the end-to-end tests and runs until it finishes, clearing 300 to 500 failing end-to-end tests a day in about four hours; people keep only the decisions that need a person.

SKILLS

LANGUAGES	TypeScript, JavaScript, Python, SQL, Bash
AI & AGENTS	Agent orchestration, LLM evals, prompt engineering, tool calling, guardrails, human-in-the-loop, context engineering, AI observability, structured output, MCP, Claude Code
WEB & BACKEND	React, Next.js, Node.js, REST APIs, OAuth, PostgreSQL
TESTING	Playwright, Appium, Page Object Model, CI test infrastructure, fault injection
INFRASTRUCTURE	Docker, GitHub Actions, CI/CD, Linux

WORK EXPERIENCE

QA Wolf

QA Lead

10/2025 to Present

- **Lead QA for Harvey, an \$11B AI legal platform:** own test strategy, the client relationship and delivery reporting; work directly with Harvey's engineers and CSMs to keep the suite running, scope new-coverage requests, and settle intended product change vs. regression from evidence
- **Manage a team of three QA engineers as a player-coach:** review every test the team submits, assign work against each engineer's bandwidth, run the team's meetings and quarterly performance reviews, and carry the same daily IC load myself
- Maintain **1,600+ Playwright/TypeScript E2E tests** across 13 product domains on an **~90-class Page Object framework (~80K lines)** with a compiler-enforced PageFactory registry; executed platform-wide POM codemod and environment migrations across the full suite
- **Handed my own job to an AI agent end to end:** it now creates, investigates, maintains and revalidates E2E tests, pulling run logs, Playwright traces, failure video and live application DOM itself, and works at the depth of a hands-on engineer, root cause and evidence attached
- **Turned around 300 to 500 failing end-to-end tests a day** in a 1,600+ test Playwright suite, usually inside a four-hour window, by running investigation and repair in parallel through the autonomous agent
- **Cut a client-facing bug report from 20-30 minutes to one keystroke:** the evidence (trace frame, log, reproduction, affected-test count) is captured during the investigation, and a second keystroke re-verifies every open report by re-running the tests behind it

QA Wolf

QA Engineer

07/2025 to 10/2025

- **Promoted to QA Lead within three months of hire** on delivery metrics
- Validated non-deterministic LLM behavior across **65 end-to-end tests:** agentic document workflows, streaming completions with multi-minute retry windows, citation integrity and knowledge-source attachment, on Microsoft Word Web and native PC, iOS/Appium, email, Twilio SMS/MFA and Google OAuth through a stealth browser

Zeta Global

Software Developer

02/2023 to 07/2025

- Built internal **React/Next.js dashboards** with real-time PostgreSQL filtering, adopted across teams as their working view into the data and retiring a standing manual reporting job
- Built a **Python collection service** over REST APIs and SQL that consolidated scattered sources into one validated pipeline, replacing manual gathering and measurably improving the accuracy of the data teams reported on
- Owned the **full deployment path:** Git, GitHub Actions and Docker pipelines that tested and deployed both systems

PROJECTS

The Control Room

Human-in-the-Loop AI Automation

Personal project, 2026 to Present

kxvin1.github.io/control-room

- Cut a broken test from **ten manual hand-offs to two:** an AI agent owns the repair, and one always-on screen shows only what needs a person
- Designed evidence gates so the agent **cannot self-certify:** "ready" needs two clean local runs plus one infrastructure run on the same commit, run id attached, with proof runs **twenty at a time; guardrails refuse destructive platform actions at the tool boundary**
- Made the agent's work **observable:** every request, reply and model call is derived from the session transcripts into an inbox with a derived status, plus a per-day **cost ledger** that measures each optimization before it is kept
- Cut agent context cost from **3-8k tokens to ~1k** per investigation with filtered evidence reports instead of raw page snapshots, and live-DOM inspection from **62s to 17s**
- Built a **measurable self-improvement loop:** every correction becomes a ledger row, a rule and a check that fails if it recurs (**600+ self-checks**, each proven against a planted failure first); repeats get their capture fixed, and the agent proposes rule changes as diffs a person accepts
- Architected panels and agent as **decoupled processes communicating only through plain-text files**, so either side restarts without losing state, and the session **hands off to itself** with a context verdict that calls a fresh session before the limit

EDUCATION

App Academy

1,000-hour full stack web development, 07/2022

CSU Northridge

B.S. Psychology, 12/2019